

Dead Code Elimination

Robbie Hammond

Case Western Reserve University,
School of Engineering
Cleveland, Ohio, USA
reh161@case.edu

Saketh Dendi

Case Western Reserve University,
School of Engineering
Cleveland, Ohio, USA
ssd74@case.edu

Milo Cassarino

Case Western Reserve University,
School of Engineering
Cleveland, Ohio, USA
mgc73@case.edu

ABSTRACT

Dead code is a big contributor for arbitrarily large program sizes and slow speed, and our project attempts to create an optimization pass that is able to eliminate easy-to-find dead code. This paper outlines some of the high-level details of our implementation of the optimization pass. In addition, this paper outlines the results of the optimization pass when run on a number of simple test cases to demonstrate its effectiveness.

KEYWORDS

compiler optimization, dead code

ACM Reference Format:

Robbie Hammond, Saketh Dendi, and Milo Cassarino. 2023. Dead Code Elimination. In *Proceedings of ACM Conference (Conference'17)*. ACM, New York, NY, USA, 3 pages. <https://doi.org/10.1145/nnnnnnn.nnnnnnn>

1 INTRODUCTION

Dead code, when unaddressed, can be a significant problem for software development. If not eliminated by the compiler, dead code can make a program larger, especially if there is a substantial amount of dead code within a codebase. In addition, dead code can make software arbitrarily slower, with additional computational resources devoted to declaring unused variables and empty functions. For these reasons, it is imperative that compilers, at a minimum, eliminate chunks of code that are unused and unneeded by the program.

Dead code, as defined in this paper, consists of two related but distinct situations. Firstly, code can be considered dead if it is never executed at runtime. For example, code inside the 'else' portion of an if-else block where the condition is always true would be dead. Secondly, a portion of code can be considered dead if the computation performed on those lines is never used anywhere else. For example, a variable that is declared but never written to or read from can be regarded as dead.

Our project focuses on adding basic dead code elimination strategies into the compiler created in the fourth class project. Our goal is to give the compiler support for eliminating unreachable code in if branches, for loops, and while loops. In addition, the compiler should be able to remove unused variables and other miscellaneous dead code optimizations.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

Conference'17, July 2017, Washington, DC, USA

© 2023 Association for Computing Machinery.

ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00

<https://doi.org/10.1145/nnnnnnn.nnnnnnn>

2 IMPLEMENTATION

In this section, we will outline how we implemented the aforementioned code removal functionality.

2.1 Class and Function Additions

Several additional functions and classes had to be created to implement our optimization pass. This section will outline the essential modifications, omitting many details and less critical additions.

2.1.1 *AST.cpp* and *function.cpp* Changes. The `Optimize()` function is the most important change to the `AST` class. The `Optimize()` function calls `PerformOptimization()` on each function. The `PerformOptimization()` function calls `CanOptimize()`, which is a function that both optimizes a statement and returns its ability to be optimized. This function is described in more detail in the subsequent sections.

2.1.2 *Statement.h* Changes. Within the `Statement.h` file, several additional functions have been declared, with the most important being `CanOptimize()` and `HowToOptimize()`. The purpose of these functions is self-explanatory. `CanOptimize()` returns whether a statement is able to be optimized. We do not try to optimize most types of statements, so the function returns false by default. `HowToOptimize()` tells the rest of the program what must be done to optimize the statement. `HowToOptimize()` returns an *Optimization*, which is an enum with values such as `REMOVE_LOOP`, `REMOVE_ELSE`, `NO_OPTIM`, and others.

2.1.3 *Block.cpp* Changes. The `CanOptimize()` function within `Block.cpp` is where much of the vital optimization work is done. As mentioned above, `CanOptimize()` is defined with `statement.h`, and its children inherit and can override this function. Though all other subclasses of `statement.h` simply return a true or false value from the function, the `CanOptimize()` function in `Block.cpp` will modify the state of the containing block class to reflect the optimization changes. Generally speaking, these state changes are the creation and deletion of statements within the statements list.

2.2 If Statements

There are two cases where a branch of an if statement will never run. The first case is when the condition is always true, regardless of the program's state when the if statement is entered. The second case is when the condition is always false.

2.2.1 *Condition Always True.* To check if the condition is always true, we get its compiled R-value. If it is 1, then we know the condition must always be true. When this happens, we don't need to have any kind of branching logic in the AST. All we need to do is generate the code within the else statement and generate the rest of the program like normal.

2.2.2 Condition Always False. To check if the condition is always false, we get the compiled R-value of the condition as in the other case. If it is 0, we know the condition must always be false. Similarly to the case when the condition is always true, all we must do is generate the code within the then statement with no branching logic.

2.3 Loops (For And While)

Just like with the if statements, there are two cases in which a loop may cause dead code. The first case is when the condition is always true, and the second case is when the condition is always false.

2.3.1 Condition Always True. Just like with if statements, we can check if the condition is always true by looking at its compiled R-value. If it is 1, then we know that the loop will never terminate, rendering the subsequent code unreachable. To simplify the AST, we can eliminate all the following statements after the loop. A return statement is still appended at the end of the loop for consistency's sake, with the return type matching that of the function. Since the code is unreachable, the value returned is unimportant, and arbitrary ones were chosen.

2.3.2 Condition Always False. If the compiled R-value of the condition is always false, there is no need to include the loop at all within the AST or LLVM IR. As such, the entire loop is erased within the AST, with the rest of the code being left unaffected.

2.4 Useless Variables

In the end this feature did not end up working as we intended although several different strategies were looked at to do so. The first was a flagging method where we would try and flag variables they were defined and only compile those that were flagged. We tried to delete variables before compilation that appeared only once in the function. We tried a method where we retroactively parsed the compiled file and removed the non-used variables from it. All of these were either wholly unsuccessful or we didn't have the time to figure them out without major problems.

2.5 Other Optimizations

An additional simple dead code elimination technique is to eliminate adjacent statements. For example, if the statement $x = 1$ is directly followed by $x = 1$, one of these statements can be removed with no effect on the code.

3 MAJOR LIMITATIONS

There still is a large number of limitations with our optimization pass. This section will outline some of the most critical ones.

3.0.1 Return/break statements within infinite loops. An infinite loop is not truly infinite if a reachable jumping statement exists within it. Our optimization is not able to recognize this, so as long as the condition is always true, it assumes the following code is unreachable.

3.0.2 Resolving more complicated conditions (specifically in regard to the AST). Our optimization pass can only recognize conditions as always being true or false when the condition is either a '0' or '1'. Obviously, this is very limiting, as most conditions that are always

true/false are more obscure. However, it is worth noting that code has been written that allows IR not to be generated for slightly less trivial conditions, such as when a variable is assigned to 0 and then used in the condition. Since the optimization pass is only supposed to operate on the AST, this logic has been commented out. See for.cpp, while.cpp, and if.cpp for how this is done.

4 TESTING METHODOLOGY

To test our optimization pass, several simple tests were created. The primary metric we decided to use to determine the effectiveness of our optimization pass was the size of the generated LLVM IR in bytes. For many of our tests, we also recorded execution time. This metric was omitted for tests containing infinite loops which last forever.

5 RESULTS

Below, you can see space and speed metrics on a small sample of test files. These test files can be found in the implementation submission. As can be seen, our optimization pass had a major impact on the resulting LLVM IR file size, with the optimized file being anywhere from 80% to 35% of the original file's size.

Also, there was not a significant impact on speed, but this is likely because the sample files tested were so short and straightforward. If a follow-up report were to be produced, we would have created a more elaborate testing suite to investigate the performance of our optimization pass on more complex pieces of code. The testing suite could include many files with a large number of loops to be eliminated and conditional branches that can be pruned.

6 CONCLUSION

Though our optimization pass has much room for improvement, it is able to recognize and eliminate obvious pieces of dead code, such as unnecessary conditional branches, useless loops, and more. Based on the testing results, we can confidently say that the removal of obvious dead code can substantially lower the size of the outputted program. In the real world, most dead code isn't so simple and easy to recognize, but our project serves as a proof-of-concept of the usefulness of dead code removal tools and optimizations.

Table 1: File Size Test Results

File	Unoptimized IR Size	Optimized IR Size
<i>if_remove_then.c</i>	848	348
<i>if_remove_else.c</i>	847	348
<i>infinite_for.c</i>	905	716
<i>useless_for.c</i>	936	321
<i>infinite_while.c</i>	888	699
<i>useless_while.c</i>	833	321

Table 2: File Execution Time Test Results

File	Unoptimized IR Size	Optimized IR Size
<i>if_remove_then.c</i>	0.009s	0.030s
<i>if_remove_else.c</i>	0.010s	0.009s
<i>useless_for.c</i>	0.009s	0.009s
<i>useless_while.c</i>	0.012s	0.009s